

Tackling Fairness, Change and Polysemy in Word Embeddings

Felipe Bravo-Marquez

Department of Computer Science, University of Chile
National Center for Artificial Intelligence Research
Millenium Institute Foundational Research on Data

July 6, 2022



dcc
CIENCIAS DE LA COMPUTACIÓN
UNIVERSIDAD DE CHILE

CEN
CENTRO NACIONAL DE INTELIGENCIA ARTIFICIAL



Millennium Institute
Foundational
Research on Data

RELELA
Representations for
Learning and Language

Word Embeddings

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

- ▶ The first step in computationally working with written language is to represent words as mathematical objects we can operate with.
- ▶ Representing words as numeric vectors a.k.a **embeddings** is a standard practice in **Natural Language Processing** (NLP).
- ▶ Word embeddings are a mapping of discrete symbols (i.e., words) to continuous vectors.
- ▶ Distance between vectors can be equated to distance between words.
- ▶ This makes easier to generalize the behavior from one word to another.
- ▶ Word embeddings have become a **core component** of NLP downstream systems (e.g., sentiment analysis, machine translation, question answering).

Word Embeddings

Tackling Fairness,
Change and
Polysemy in Word
Embeddings


Felipe Bravo
Márquez

Introduction


Fairness

Change


Polysemy



Male-Female



Verb tense



Country-Capital

Word embeddings can encode semantic and syntactic relationships between words.

Distributional Hypothesis

- ▶ The construction of word embeddings from document corpora is based on the **Distributional Hypothesis** [Harris, 1954]:
*Words occurring in the same **contexts** tend to have similar meanings.*
- ▶ Or equivalently:
*A word is characterized by the **company** it keeps.*
- ▶ The idea of word embedding is to map words occurring in similar contexts to similar vectors.

Word Embeddings

Tackling Fairness,
Change and
Polysemy in Word
Embeddings


Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy



Word Embeddings Algorithms

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

- ▶ Word embeddings are built by training **neural networks** architectures on document **corpora** (e.g., books, papers, Wikipedia, tweets, the Web).
- ▶ These architectures formulate a **predictive task** (e.g., predict a missing word within a context window) in which word embeddings naturally arise from the network's parameters after training.
- ▶ Most popular models are:
 - ▶ Skip-gram negative sampling [Mikolov et al., 2013]
 - ▶ Continuous bag-of-words [Mikolov et al., 2013]
 - ▶ Glove [Pennington et al., 2014]
 - ▶ FastText [Bojanowski et al., 2016].

Skip-gram Model

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

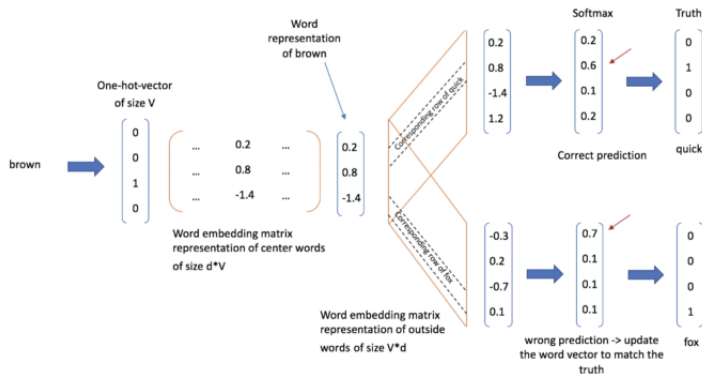
Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy



⁰Picture taken from: <https://towardsdatascience.com/word-embeddings-intuition-and-some-maths-to-understand-end-to-end>

Limitations of Word Embeddings

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

- ▶ In this talk we will present our research addressing three limitations of word embeddings.
 1. **Fairness**: they are prone to inherit stereotypical social biases from the corpus they were built on.
 2. **Change**: they are static. Thus they ignore words not observed during training and are unable to capture semantic drifts.
 3. **Polysemy**: they fail to capture the polysemous nature of many words (e.g., apple:company, apple:fruit), conflating their multiple senses into a single point.

Fairness in Word Embeddings

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

Word embeddings have shown to exhibit **undesirable associations** for certain social identity groups (e.g., gender, race, religion):

- ▶ Black is to criminal as caucasian is to police.
- ▶ Man is to doctor as woman is to nurse.
- ▶ Man is to computer programmer as woman is to homemaker.

- ▶ **WEAT** [Caliskan et al., 2017]: measures the degree of association between two sets of target words and two sets of attribute words.
 - ▶ The original paper focused on biases regarding ethnicity (in relation to pleasant-unpleasant words) and gender (in relation to occupations).
- ▶ **RND** [Garg et al., 2018]: compares embeddings related to two different social groups against a set of neutral words.
 - ▶ The paper evaluated gender and ethnic biases in embeddings trained over different time periods.
- ▶ **RNSB** [Sweeney and Najafian, 2019]: measures the KL divergence between the negative sentiment probability of the embeddings and a uniform distribution.
 - ▶ The paper studied bias towards national origin identity terms such as American, Mexican, and Canadian.

WEFE: The Word Embeddings Fairness Evaluation Framework

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

- ▶ Although metrics have similar objectives they cannot be easily compared with each other.
- ▶ They operate with different inputs (e.g., pairs of words, sets of words, multiple sets of words).
- ▶ Their outputs are also incompatible with each other (e.g., reals, positive numbers, $[0, 1]$ range).
- ▶ The **Word Embeddings Fairness Evaluation (WEFE)** is a framework for measuring and mitigating bias in word embeddings [Badilla et al., 2020].
- ▶ WEFE formalizes existing fairness metrics into a **unified framework**.

WEFE: The Word Embeddings Fairness Evaluation Framework

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

- ▶ It encapsulates the input words of these metrics into standard objects called **queries**.
- ▶ Then, proceeds to quantify the bias of word embeddings using these queries along with different fairness metrics.
- ▶ It also allows us to **rank** various embedding models by aggregating the previous scores.



WEFE Building Blocks

- **Target Sets:** a set of words intended to denote a particular social group.
For example:

$$T_{\text{women}} = \{w_{\text{she}}, w_{\text{woman}}, w_{\text{girl}}, \dots\}, \quad (1)$$

$$T_{\text{men}} = \{w_{\text{he}}, w_{\text{man}}, w_{\text{boy}}, \dots\}, \quad (2)$$

- **Attribute Sets:** a set of words representing some attitude, characteristic, trait, occupational field, etc. that can be associated with individuals from any social group.
For example:

$$A_{\text{science}} = \{w_{\text{math}}, w_{\text{physics}}, w_{\text{chemistry}}, \dots\}, \quad (3)$$

$$A_{\text{art}} = \{w_{\text{poetry}}, w_{\text{dance}}, w_{\text{literature}}, \dots\}. \quad (4)$$

- **Query:** a pair $Q = (\mathcal{T}, \mathcal{A})$ in which \mathcal{T} is a set of target word sets, and \mathcal{A} is a set of attribute word sets.
For example:

$$Q = (\{T_{\text{women}}, T_{\text{men}}\}, \{A_{\text{science}}, A_{\text{art}}\}). \quad (5)$$

- ▶ Queries are the main building blocks used by fairness metrics to measure bias of word embedding models.
- ▶ Fairness metrics encapsulated by WEFE are defined with a query template $s_F = (t, a)$, which is a pair of natural values with the cardinalities of sets \mathcal{T} and \mathcal{A} .
- ▶ We equip each metric with a total order relation \leq_F that establishes what is to be considered as *less biased*.
- ▶ That is, if

$$F(\mathbf{M}_1, Q) \leq_F F(\mathbf{M}_2, Q)$$

we can say that \mathbf{M}_1 is *less biased than model* \mathbf{M}_2 .

- ▶ With this order relation we can rank various embeddings models according to their biases.

WEFE Toolkit

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

We released WEFE as an open source Python software:

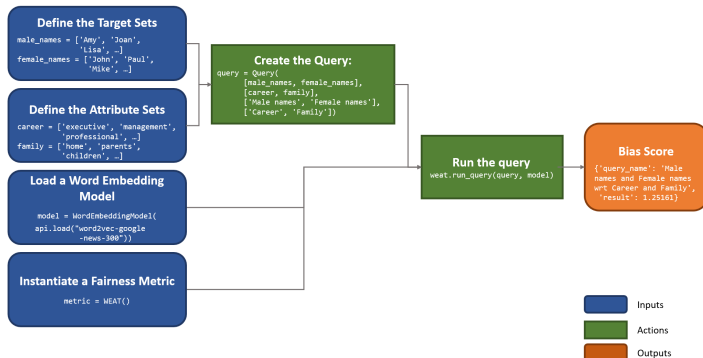
<https://wefe.readthedocs.io/>

Introduction

Fairness


Change

Polysemy



Case Study

We used WEFE to compare various embedding models according to their aggregated biased with respect to gender, religion and ethnicity:



Cumulative ranking according to overall metric scores

Incremental Word Vectors for time-evolving sentiment lexicon induction

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

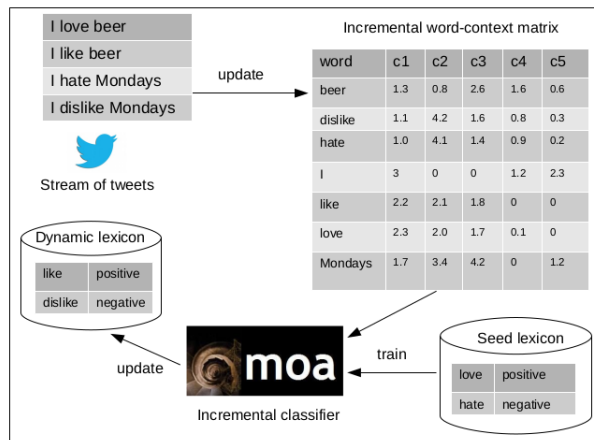
Polysemy

- ▶ A sentiment lexicon is a list of expressions annotated according to affect categories such as positive, negative, anger and fear.
- ▶ Lexicons are widely used in sentiment classification of tweets.
- ▶ They are usually build by training a word-level sentiment classifier on pre-trained word embeddings using a **seed lexicon** as training data [Tang et al., 2014].
- ▶ However, as word embeddings are static by nature, sentiment lexicons build in this way are unable to capture:
 1. The arrival of new sentiment-conveying expressions such as #trumpwall and #PrayForParis.
 2. Temporal changes in sentiment patterns of words (e.g., a scandal associated with an entity).

Incremental Word Vectors for time-evolving sentiment lexicon induction

We developed a methodology for automatically inducing **continuously updated** sentiment lexicons from **Twitter streams** by training **incremental** word sentiment classifiers from **time-evolving word vectors**.

[Bravo-Marquez et al., 2021]



Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

Incremental Word Vectors for time-evolving sentiment lexicon induction

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

The process is summarized as follows:

1. Connect to a stream of continuously arriving text (e.g., Twitter).
2. Every time a new tweet arrives, a sliding window of W words centered on a target word is shifted across the content.
3. The word vector associated with the target word is updated according to its context. This process implies updating word-context counts and computing PPMI-based associations.
4. If the target word is new, a new vector associated with this word is created.
5. If the target word is contained in the seed lexicon, the incremental classifier is updated/trained using the target word vector and the lexicon's sentiment as the gold label.
6. If the target word is not contained in the lexicon, its sentiment is estimated using the classifier and the dynamic lexicon is updated.

Introduction

Fairness

Change

Polysemy

Incremental Algorithms

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

Data: tweets, window size W , vocabulary size V ,
context size C

Result: PPMI Matrix M

Initialize PPMI Matrix M of size $V \times C$;

$D \leftarrow 0$;

while *tweet in Data Stream* **do**

 tokens \leftarrow tokenize(tweet);

for each w in tokens **do**

$D \leftarrow D + 1$;

$c_1, \dots, c_{2W} \leftarrow$ getContexts(w , tokens);

 updateContextMatrix(w , c_1, \dots, c_{2W});

for each c_j in c_1, \dots, c_{2W} **do**

 PPMI(w , c_j) \leftarrow

$\max \left(0, \log_2 \left(\frac{\text{count}(w, c_j) \times D}{\text{count}(w) \times \text{count}(c_j)} \right) \right)$;

end

end

end

Data: tweets, training seed lexicon, testing seed
lexicon

while *Tweet T in Data Stream* **do**

for each token t in tokens **do**

 updateContextCounters(t , T);

 updatePPMIValues(t , T);

if token \in train seed lexicon **then**

 trainClassifier(PPMI(t), label(t));

else if token \in test seed lexicon **then**


 updateEvaluator(PPMI(t), label(t));

end


end

Sentiment Drifts

- We simulate sentiment change by randomly picking some words and swapping their context with the context of words exhibiting the opposite sentiment.



1. Accuracy for experiment with 30% of Change

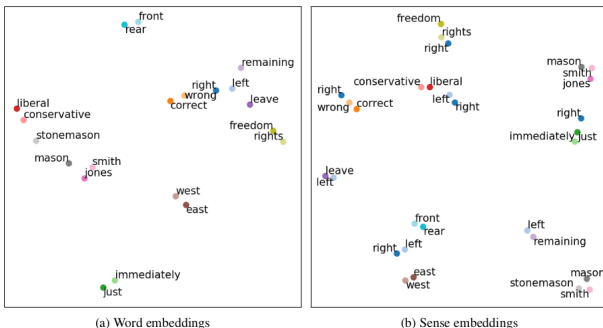


1. (a) dead

- Our approach allows for successfully tracking of the sentiment of words over time even when drastic change is induced.

PolyLM: a polysemous language model

- PolyLM is a language model capable of automatically learning multiple meanings of a word (e.g. apple:apple, apple:company) [Ansell et al., 2021].



- The occurrence of closely related polysemous words nearby in the word embedding space (i.e. *left* and *right*) causes unrelated words to be closer together (e.g. *left* and *wrong*) and related words to be further apart (e.g. *right* and *east*) than they otherwise would be

PolyLM: a polysemous language model

- ▶ PolyLM is an unsupervised language model based on two assumptions:

1. The probability of a word occurring in a given context is equal to the sum of the probabilities of its individual senses.

$$p(w|c) = \sum_{s \in S_w} p(s|c). \quad (6)$$

2. For a given occurrence of a word, one of its senses tends to be much more plausible in the context than the others.

- ▶ PolyLM is based on a Transformer encoder architecture and the masked language modeling (MLM) task used for training BERT.
- ▶ Its loss function combines a standard language modeling loss with a distinctness loss encouraging one sense to have a high estimated probability of occurring, and the rest to have a low probability.

PolyLM: a polysemous language model

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

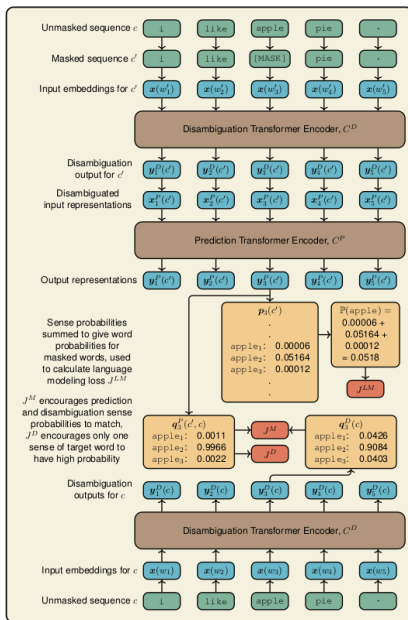
Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy



Results on the Word Sense Induction (WSI) task

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

System	Version	SemEval-2010			SemEval-2013		
		F-S	V-M	AVG	FBC	FNMI	AVG
Aamrami-goldberg AutoSense	BERT _{LARGE}	71.3	40.4	53.6	64.0	21.4	37.0
		62.9	10.1	25.2	61.7	8.0	22.2
PolyLM [†]	BASE	65.8	40.5	51.6	64.8	23.0	38.6
	SMALL	65.6	35.7	48.4	64.5	18.5	34.5
Qiu et. al [†]		-	-	-	56.9	6.7	19.5
SE-WSI-fix-cmp [†]		54.3	16.3	29.8	-	-	-
AdaGram [†]		43.9	20.0	29.6	13.2	8.9	10.8
Arora et. al [†]	$k = 5$	46.4	11.5	23.1	-	-	-

Table: Comparison of sense embedding models and WSI-specific techniques on the SemEval 2010 and 2013 WSI tasks. SE-WSI-fix-cmp is based on MSSG model. [†] - models which obtain explicit sense embeddings.

Conclusions

- ▶ In this talk we have presented three research projects, each addressing a different limitation of word embeddings.
- ▶ I would like to thank all the co-authors of the presented works: Alan, Arun, Jorge, Pablo and Bernhard.
- ▶ There is plenty of room for further research.

Questions?

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy

Thanks for your Attention!

References I



Ansell, A., Bravo-Marquez, F., and Pfahringer, B. (2021).
Polym: Learning about polysemy through language modeling.
In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume, pages 563–574.



Badilla, P., Bravo-Marquez, F., and Pérez, J. (2020).
WeFe: The word embeddings fairness evaluation framework.
In IJCAI, pages 430–436.



Bojanowski, P., Grave, E., Joulin, A., and Mikolov, T. (2016).
Enriching word vectors with subword information.
arXiv preprint arXiv:1607.04606.



Bravo-Marquez, F., Khanchandani, A., and Pfahringer, B. (2021).
Incremental word vectors for time-evolving sentiment lexicon induction.
Cognitive Computation, pages 1–17.



Caliskan, A., Bryson, J. J., and Narayanan, A. (2017).
Semantics derived automatically from language corpora contain
human-like biases.
Science, 356(6334):183–186.



Garg, N., Schiebinger, L., Jurafsky, D., and Zou, J. (2018).
Word embeddings quantify 100 years of gender and ethnic stereotypes.
Proceedings of the National Academy of Sciences,
115(16):E3635–E3644.

References II



Harris, Z. (1954).
Distributional structure.
Word, 10(23):146–162.



Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013).
Distributed representations of words and phrases and their
compositionality.
In Burges, C., Bottou, L., Welling, M., Ghahramani, Z., and Weinberger,
K., editors, *Advances in Neural Information Processing Systems 26*,
pages 3111–3119. Curran Associates, Inc.



Pennington, J., Socher, R., and Manning, C. D. (2014).
Glove: Global vectors for word representation.
In *Proceedings of the 2014 Conference on Empirical Methods in Natural
Language Processing, EMNLP 2014, October 25-29, 2014, Doha, Qatar,
A meeting of SIGDAT, a Special Interest Group of the ACL*, pages
1532–1543.



Sweeney, C. and Najafian, M. (2019).
A transparent framework for evaluating unintended demographic bias in
word embeddings.
In *Proceedings of the 57th Annual Meeting of the Association for
Computational Linguistics*, pages 1662–1667.

References III

Tackling Fairness,
Change and
Polysemy in Word
Embeddings

Felipe Bravo
Márquez

Introduction

Fairness

Change

Polysemy



Tang, D., Wei, F., Qin, B., Zhou, M., and Liu, T. (2014).

Building large-scale twitter-specific sentiment lexicon : A representation learning approach.

In Proceedings of the 25th International Conference on Computational Linguisticss, pages 172–182. Association for Computational Linguistics.