

Introduction to Reinforcement Learning

Felipe José Bravo Márquez

December 30, 2022

Markov Decision Process

- A markov decision process is a tuple:

$$(S, A, \{P_{SA}\}, \gamma, R)$$

where

- S is a set of states.
- A is a set of actions.
- P_{SA} are the state transition probabilities:

References I



Bickel, P. J., Hammel, E. A., and O'Connell, J. W. (1975).

Sex bias in graduate admissions: Data from berkeley: Measuring bias is harder than is usually assumed, and the evidence is sometimes contrary to expectation. *Science*, 187(4175):398–404.



Hardt, M. and Recht, B. (2021).

Patterns, predictions, and actions: A story about machine learning. *arXiv preprint arXiv:2102.05242*.